

ALife Models as Epistemic Artefacts

Xabier Barandiaran¹ and Alvaro Moreno¹

¹ University of the Basque Country
Department of Logic and Philosophy of Science, FICE,
UPV-EHU (University of the Basque Country)
PO BOX 1249 / 20080 Donostia - San Sebastian / Spain
xabier@barandiaran.net
<http://www.ehu.es/ias-research>

Abstract

Both the irreducible complexity of biological phenomena and the aim of a universalized biology (life-as-it-could-be) have lead to a deep methodological shift in the study of life; represented by the appearance of ALife, with its claim that computational modelling is the main tool for studying the general principles of biological phenomenology. However this methodological shift implies important questions concerning the aesthetic, engineering and specially the epistemological status of computational models in scientific research: halfway between the well established categories of theory and experiment. ALife models become powerful epistemic artefacts allowing the simulation of emergent phenomena, the interaction between different levels of organization and the integration of different causal factors in the very same manipulable object. The use of computational models in ALife can be classified in four main categories depending on their position between theoretical and empirical practices: generic, conceptual, functional and mechanistic. For each of these categories we analyse their epistemic value and select paradigmatic examples that illustrate how ALife models can be fruitfully inserted in the study of life.

Introduction

Modern Biology shows living systems as highly complex networks of many interrelated elements structured in different levels, which in turn interact among them at different time-scales and hierarchies, where local contingencies and universal self-organizing properties merge to produce a particular phenomena. Unlike other natural phenomena, such complex organizations are not describable as the action of few simple laws because many local rules, appearing at different levels, constrain the dynamical processes in intricate ways. In addition, for this kind of systems traditional analytic and experimental research strategies are highly limited. On the analytic side, decomposition of a system into functionally complete structures is often impossible, since functionality and global structural patterns emerge out of non-linear interactions between lower level components and system-environment recursive interactions. On the experimental side, intensive control and manipulation of the studied system is not always possible.

In order to study such systems, a synthetic approach is required in which biological patterns can be reproduced according to the abstraction of the underlying mechanisms that produce them. As E. Fox Keller (2005) has pointed out, the complexity of biological systems can only be modelled through vast systems of differential equations, statistical analysis, larger sets of algorithms and simulations, whose partial results require an intimate back and forth relation with empirical experiments. Given this, the study of living systems requires special techniques, many of which have been made possible due to the development of computers. Unlike the limitations of traditional mathematical models of theoretical biology, where models must be simplified to accommodate analytic tractability, computer models permit a much wider range of formal treatment by means of numerical methods. At the same time, automatized experimentation in the computer allows a systematic study of the effect of different parameter values, perturbations on the system, change of initial and boundary conditions, and a host of intensive explorations in the model. All this range of manipulations cannot be performed on the natural system under investigation.

Hand by hand with a quantitative and systematic study of different biological organizations (ecosystems, metabolic, genetic, neural networks, etc.) it appears also the possibility to go further than the empirical instances corresponding to the historical contingencies of terrestrial evolution. Based on the computational “experimentation”, it would be possible to formulate the organizational principles of all possible life, including unknown natural forms of life and even new artificially created ones. In other words, these new methodological tools have opened the possibility to universalize Biology. This picture of biological organization (as well as its computational synthetic and universalizable investigation) is what ALife brings forth through a set of simulation techniques and artefacts that have been developed over the years, constituting a new kind of science: the study of life-as-it-could-be (Langton 1989).

The price to pay for universalization and computer experimentation is the absence or distance from a direct reference to any explicit empirical domain. Thus, an adequate interpretation of the functional meaning of the

different processes and patterns, obtained when running the model, becomes much more problematic than what traditional epistemology of scientific methodology has established. ALife simulations acquire a hybrid status; they are, as Sismondo (1999) has pointed out, at the same time tools, representations, objects and ideas.

Unlike other disciplines or fields, within the ALife community simulation models sit particularly comfortable. They are not perceived as threatening the methodological identity of the field. On the contrary, simulation models constitute a unifying factor of the community which is, itself, a transdisciplinary field whose fuzzy boundaries are defined by an umbrella of techniques; transcending themselves into a philosophy of biology, cognition and complex systems in general. The experimental side of the ALife community does not only lie on the results of its simulation models, but on a continuous exploration of new modelling techniques. As a consequence, traditional paradigms of epistemology, methodology and even ontology are always being challenged. Thus, although socially accepted, the epistemic status of simulation models has long been a controversial issue within the community itself (Wheeler et al. 2002, Webb 2001, Di Paolo et al 2000, Moreno 2000, Bonabeau & Theraulaz 1994, Pattee 1989).

This paper is not intended to be what Langton (1994) termed a YAPP (Yet Another Philosophical Paradigm), but to provide a conceptual guide which makes explicit the different goals that an ALife model may pursue and to classify their evaluative frameworks, which becomes a necessary condition to avoid the opposite risk, i.e. what Langton calls YANN (Yet Another Neural Net) style papers. Although most of the major contributions within the field somehow reinvents its epistemological status, some general considerations can be made about the epistemic use and value of ALife models (ALM hereafter). In this paper we distinguish three main modalities of ALMs: aesthetic, engineering and epistemic. Then, we classify ALM in four levels of abstraction specifying an epistemic evaluative framework for each of them.

Simulations, realizations, theories and models in ALife

Although in this paper we will focus on the study of computer simulations in ALife, it is important to clarify previously the relations that theories and models have with simulations and also realizations.

A theory is a structure of concepts in the human mind aimed to explain a given domain of reality. This structure of concepts gives rise to simulations and/or realizations, through the articulation of these concepts in a model. Simulations are formal, rate-independent, unfolding of mathematical models in computer media. As Pattee (1989) has pointed out, in a simulation a set of formal rules try to reproduce the laws and material constraints that govern the

dynamics of certain physical systems, or even, the simulation may generate a virtual universe of objects following alternative laws. A realization, on the contrary, is a physical construction of a system to be used as a model for the study of a given phenomenon. A realization is an example of the synthetic methodology of ALife, since it seeks the understanding by a literal reproduction of those elements supposedly fundamental in the generation of a given phenomenon. When the model gives rise to a realization, the processes that unfold the consequences of the basic conceptual assumptions are (at least in part) real, rate-dependent, dynamical processes, so that the system literally obeys the physical laws.

An important proviso is worth taking in mind: in a certain sense, a model is more than the simulation or realization of a system, since it includes auxiliary assumptions to interpret the computational or physical system as a relevant tool. In each case, there is a host of auxiliary knowledge, information, assumptions or transformations, which is never explicit in the artificial construct itself, but becomes an essential component of the model. We can imagine the model to be a physical or computational construct (realization or simulation) plus an envelop of auxiliary knowledge and techniques that permit to use the construct as a tool that somehow captures the causal structure of a given phenomena or theoretical domain. Failure on making such construct productive might be due to failure on this auxiliary framework, that makes it a valuable or non-valuable instrument. Such a framework is as important as the programming elements of the simulation in the success of a model. But while computational systems have standard methods of testing and debugging the auxiliary framework is not always easy to evaluate itself, it depends on a set of, not always explicit, assumptions. In fact, as we shall see, the epistemic value of many ALMs (especially of those we shall call conceptual) lies on their capacity to make explicit, dissonant or inconsistent some of the components of these auxiliary frameworks, thus leading to reorganization of some theoretical issues.

Aesthetic, Engineering and Epistemic goals for ALife Models

One of the main characteristic of ALife modelling techniques is their capacity to produce emergent patterns or behaviours out of a set of local rules or equations. These emergent patterns can be used for different purposes. For instance, models can be used as *epistemic* tools to understand how natural systems work, to explain life and living phenomena, making explicit and tractable the underlying universal principles of biological organization by exploring properties of complex systems and simulating some specific natural behaviour. But that is not the only use an ALM can certainly have. The emergent order that natural systems show can also be seen as the unfolding of a functional evolutionary design that turns out to be

attractive for *engineering* goals. In fact, evolution is an extremely accurate blind engineer that can teach us weird and new ways of achieving solutions for a wide range of problems. Both the organization of individual living systems and their collective behaviour (e.g. colonies) also give rise to intriguing designs and functional behaviours. Nature is not constrained by any theoretical assumption, design principles or industrial tradition, and thus becomes an attractive pool, in which original and sophisticated, simple but powerful engineering techniques can be found. Finally, under *aesthetic* purposes biologically inspired artefacts can also be used to produce pleasant and original visual or musical patterns.

In each case the aesthetic, engineering or epistemic use of ALMs requires that the model be evaluated, interpreted and designed in different (although often overlapping) terms, i.e. the modelling framework of the artificial system will depend on these three main kind of goals. For instance, artists will evaluate ALMs against the aesthetic quality of the patterns they produce, their resemblance with some natural aesthetic principles, or the models' capacity to generate visual or musical patterns for an artist to select and combine in a performance. When used for engineering purposes, ALMs and artefacts are tested or evaluated against a desired functional outcome and adapted to a set of material or computational resources. Epistemic purposes involve a much more complex evaluation framework, ranging from a realistic correspondence with a particular organism or biological process to the exploration of highly abstract properties of complex systems. It is this final purpose what shall concern us the most throughout the rest of this paper. We want to understand how an artificial system becomes an epistemic artefact, what kind of knowledge it can produce, and how these artefacts interact with theoretical and empirical objects.

Often, in ALife, the temptation of building artificial worlds capable of generating some aesthetically pleasant pattern, presenting a vague biological resemblance and based on a partially new computational technique, is too high; but the resulting model might easily end up being equally useless for any of the previously mentioned purposes. That is why it is important to define an specific purpose for the model in order to adequately establish a framework for its evaluation and, hence, for its design, interpretation and experimentation. In some cases, though, the boundaries between the proposed three categories might not be very clear. For instance, engineering and epistemic purposes might be intertwined when the level of abstraction of the model is so high and generalist that it covers both the actual (a wide range of existing natural phenomena) and the possible (and thus useful for engineering purposes). On the other hand, "wet" ALMs aiming to produce minimal living cells easily cross the boundaries between engineering and epistemic purposes. In fact, the universalist aspiration of ALife research, its foundational focus on life-as-it-could-be rather than life-as-we-know-it (Langton 1989) and the functional nature of

living system threatens the boundary between science and engineering. In addition, most ALife research feeds back to itself, ALife is certainly a technoscientific field where technological artefact and objects of study merge.

But, even in limit cases, the boundaries are worth keeping. Within ALife there are still significant differences between engineering and epistemic purposes that make a difference on the construction, testing and interpretation of models. On the one hand, when talking about the naturally possible and the artificially possible, different constraints apply. Let us take the example of an artificial system where some kind of self-assembly occurs. If this system is used as a model of natural self-assembly (epistemic use) it will have to satisfy completely different constraints than if used as a model for artificial self-assembly (engineering). For instance in the first case the model should accommodate natural thermodynamic constraints. Whereas the engineering use of the system could lead to a model where the self-assembling components have access to global information obtained from a panoramic camera or where energy requirements are solved using batteries. What would be misleading in this case would be to generalize claims on natural self-assembly on the basis of the engineering model or claiming engineering limits based on a number of assumptions derived from spontaneous natural processes. On the other hand, the experiments carried out with the model might be very different if we are to highlight one or the other purpose: it is not the same to make experiments in order to find the relevant parameters to produce harmonic patterns (aesthetic purpose) than experimenting to maximize energy efficiency (engineering purpose). Finally the way to contextualize and show the experimental results will depend on the chosen framework. For instance for engineering purposes we might want to compare the proposed model with other existing techniques to solve a class of problems, while in the epistemic case we might want to highlight the plausibility of the model with, lets say, what the chances are that the modelled phenomena could appear under the environmental conditions we can expect to happen in other planets.

From idealism to empiricism: 4 epistemic uses of ALMs

The main interest and methodological novelty of ALMs, from the point of view of science and epistemology, lies in its capacity to develop an experimental research program in the computational domain. But once we fix an epistemic purpose for a model, there are still many different frameworks for its evaluation. The epistemic status of this computational domain does not fit within traditional concepts of philosophy of science. There are no standardized methods to achieve fruitful scientific results from ALMs. This condition requires that, together with the model, we carefully design and specify how it shall be evaluated. For scientific-epistemic purposes we can

distinguish four different classes of such frameworks, according to the level of abstraction the model shall adjust to or, in other words, the position models are made to occupy between theories and empirical data. This classification in four levels is certainly not exhaustive, but reflects somehow a clustering of existing models and their underlying epistemic logic.

We distinguish those models whose object of modelling is theoretical (whose evaluation is done against theoretical principles or formalisms), and those whose object is empirical (which are evaluated against empirical data coming from specific natural phenomena). In the first class we distinguish, on the one hand, those generic systems that serve to discover or classify generic properties of complex systems; we call these *generic models*. On the other side of the theoretical realm, we find what we call *conceptual models*, whose framework of evaluation is given by the dissonance or adequacy with the relationships established among concepts of a given theory. Abstract conceptual models are used to formalize or compare definitions of generic concepts (such as emergence, complexity or hierarchy) while domain specific conceptual models are used to explore the role and interaction between more specific concepts (such as learning, autopoiesis, cognitive agent, evolvability, etc.). The third and fourth classes of epistemic models directly involve their matching or interaction with empirical data. On the one hand, *functional models* are defined as those that must adjust to the particular behaviour or functionality exhibited by certain natural system. *Mechanistic models*, on the other hand, are those which act as functional and structural models of a particular phenomena. The model is meant to be realistic: variables of the computational system represent observables of the empirical system, including the internal mechanisms producing the phenomena under study.

Note that the same artificial system can be used at these four different levels of abstraction, so that it is not the simulation itself what specifies its level of abstraction, but the use and interpretation of it. For instance CTRNNs (Continuous Time Recurrent Neural Networks) can be used at different levels: to explore abstract systemic properties of dynamic systems at an abstract generic level, to explore the concepts of learning and memory at a conceptual level, to achieve a desired functionality found in ant behaviour at a functional level, or to model the pattern generator circuit of Aplysia at a more realistic mechanical level.

Generic Models

At the most abstract level we find certain computational constructions with no particular reference to any specific object of study, but whose formal structure has been selected in virtue of its resemblance with a wide range of natural phenomena. Their generality is very high and exploration of these artificial systems leads to the discovery of generic abstract properties of complex

systems. Such is the case of explorations in Random Boolean Networks, Cellular Automata (CA), Scale-Free Networks, some research in Neural Networks, Dynamical Systems, etc. Often generic models evolve out of empirical or functional models when a particular structure is found to have properties which might be generalizable to other domains; then particular and domain specific details of the model are abstracted and the empirical model is transformed into a generic one.

A very common methodology to extract valuable knowledge from these models is the exhaustive statistical analysis of different configurations of the system as well as measurements and operations performed on the resulting patterns. As a consequence, the space of resulting behaviours or structures might be classified and a set of internal relationships between parameters and resulting patterns are discovered. This knowledge of the properties of generic models can feed back to the empirical domain both at predictive and constructive levels, given the appropriate adaptation of the abstract system to the real one.

A paradigmatic example of generic systems is provided by scale-free networks. Although the theory of scale-free networks was originally developed by physicist analysing the structure of the WWW (Barabási and Reka 1999, Barabási 2002) it has gained increasing attention in ALife. Research in scale-free networks has discovered a number of properties of the structure of non-random networks found in nature: small-world phenomena (any two nodes are separated by a very small number of connections), robustness to random node failures, etc. The same properties and structure apply to social networks, protein-interaction networks, computer networks, economic networks, etc. Stuart Kauffman's explorations into generic properties of Random Boolean Networks (Kauffman 1969, 1993, Gershenson 2004) are another instance of generic models aiming to discover some universal principles of self-organization in complex systems (independently of the social, genetic, metabolic, immune or neural nature of these systems).

Conceptual models

The early enthusiasm within ALife of the genuine instantiation of living phenomena in the computer gave rise to a whole a set of artificial worlds. Discussion about the plausibility, similarity, and adequacy of these models with existing theories of life produced both a fruitful debate about the underlying assumptions in theoretical biology and the status of such models as instances of living beings. As a result, models moved to be used as a tool to question and reorganize theoretical assumptions and concepts rather than to the creation of so claimed artificial living systems without a clear epistemic purpose. Some of these models became conceptual ones which are, probably, ALife's most specific and original use of models.

What we call a conceptual model involves the simulation of processes which are, in virtue of some dynamic or structural analogy with theoretical notions, *conceptualized* under a certain theory of the living, cognitive, social or, in general, complex systems. Conceptual models can be very abstract or very specific depending on the theory under which they are interpreted/constructed. For instance, at the abstract level, the model could work to illustrate, formalize or compare one or more theories of emergence using, let's say, CA patterns. On the other hand, a domain specific conceptual model can be exemplified by a simulation of active perception in situated agents.

Unlike generic models (whose applicability to the empirical realm is more standard due to its similarity with abstract mathematical research) conceptual models have a more heterodox epistemic status: their relation to theories and empirical data is more complex and intricate. As C. Emmeche (1994) has pointed out, these models (which he calls "second-order simulacra") are not elaborated as abstractions of the biological empirical domain, but from the biological theories themselves. Such kind of models are not attempts to capture some empirical reality but are abstract worlds in their own right (e.g. Tierra). Maynard Smith has called AL into question arguing that it is a "science without facts", referring to the problem of how to assess a set of computational models whose (potential) empirical references are imprecise or non-existent (quoted in Horgan 1995). However, it would be an error to evaluate conceptual ALM based research in these terms. The main interest (and methodological novelty) of conceptual ALM lies probably in its capacity to develop experimental research on the internal conceptual relationships within theories of biological organization. This computational research allows what Dennett (1994) calls the realization of highly rigorous and far-reaching thought experiments, which the naked human mind never could perform. Bedau (1998) and, particularly, Di Paolo and colleagues (Di Paolo et al 2000) have elaborated a more detailed account of the role and methodology of ALMs as "opaque thought experiments". The opacity of the thought experiment lies on the complexity of the model. The unfolding of properties and patterns from a set of premises (local rules or differential equations) are not always predictable in the absence of a computer simulation that performs recursive calculations, integrates random perturbations, visualizes the results and so on. As it happens with traditional thought experiments, the epistemic value of conceptual ALMs does not lie on their adequacy with some empirical phenomena (since the thought experiment involves hypothetical and idealised situations). On the contrary, the model operates on the hidden assumptions of the theories used to design and interpret the model and on the conceptual relationships between these assumptions.

When concepts of a theory are related to each other through relationships which cannot always be derived on logical grounds, computer simulations become cognitive

tools for theoretical development (Casti 1997). For instance learning and ontogenetic plasticity has intricate effects on evolution. The interaction between these two concepts (learning and evolution) is difficult to generalize and study through natural fossil records or other empirical means. An alternative is to develop artificial worlds (whose local rules are abstractions of the generic mechanisms that evolutionary theory takes to be essential for natural evolution) where simplified forms of evolution and learning can be studied. The Baldwin effect (for example) was nicely illustrated by a computer model (Hinton and Nowlan 1987). Subsequent ALife research has made explicit many other properties and dynamic relationships between learning and evolution (Ackley and Littman 1992, Suzuki and Arita 2004).

The conceptual relationships that the models uncover, illustrate or deny are rarely the result of just running the simulations. Most of these models require a careful exploration and experimentation under different conditions in order to generalize results and find intermediate explanatory patterns to extract conceptually useful knowledge from the dynamics of the simulation model. The resulting conceptual achievement shall later be used to configure explanatory patterns of specific cognitive, evolutionary, metabolic or collective models subject to empirical manipulation and introduced on the traditional scientific hypothetico-deductive method.

Conceptual ALMs are used in a number of stereotyped ways. Proofs of concept are a use of these models in which the possibility to produce a particular behaviour is demonstrated in the model, given a set of mechanisms previously considered incapable to produce such behaviour or functionality (good examples of proof of concept are Seth 1996 and Beer 2003). On the other hand models are often used to illustrate, formalize or quantify a previously ill-defined concept such as emergence, hierarchy, autopoiesis, representation, etc. Here the model acts as a simplified environment in which (given that the competing theoreticians accept the assumptions on which the model is based) disputes over a conceptual definition can be solved in virtue of their viability, accuracy, correlation with expected classification, applicability to known phenomena, etc. Conceptual models are also used to study interactions happening between different levels of organization or different scales, which are often studied by different disciplines, with different theories and tools. Examples of these levels that appear merged in the same simulation are neural mechanism and behaviour, genotype and phenotype, ontogenetic development and evolution, individual and collective behaviour, etc.

Functional models

Traditional explanatory strategies are based on functional/task decomposition and structural localization of such functions so that the functioning of a system can be understood by the aggregation of such component parts (Bechtel and Richardson 1993). Complex systems, on the

contrary, are not amenable to this cognitive strategy, they are integrated systems where the overall functionality of the system emerges out of internal recursive interactions between components and system-environment interactions. That is why functional models become necessary tools to study a complex integrated behaviour. These models are tested against a particular behaviour or functionality exhibited by some natural system, but where the mechanisms that produce this functionality in the natural process are unknown, controversial or incompletely understood. Unlike conceptual models, empirical functional models must include constraints which are specific of the phenomena under investigation. These models are used to discover candidate mechanisms or local rules that produce or contribute to the observed and simulated global pattern or behaviour, to assess the performance of existing hypothesized candidate mechanisms, to assess which environmental factors participate in the causal structure of the behaviour, etc. A paradigmatic example of functional models is found among research on swarm-intelligence (Deneubourg et al 1991, Liviú and Luke 2004), where colonies' foraging capacities are reproduced and studied as a result of stigmergic and cooperative interactions.

Particularly interesting instances of functional models are given by the use of CTRNNs (and its variants) as universal smooth dynamical system approximators (independently of any attributed resemblance with real neuronal architectures). Artificial evolution is applied to a CTRNN control system to achieve a particular embodied behaviour on a simulated robot (Beer 2003, Vickerstaff and Di Paolo 2005). The dynamic causal structure of the resulting behaviour can then be analysed independently of the natural mechanism that could support it. Halfway between conceptual and functional modelling this research strategy leads to a kind of emergent functionalism in which it is the dynamic organization of behaviour what captures the essential features of cognitive processes as opposed to the representational properties of symbols operated by propositional rules (as in cognitivist or computational functionalism).

Mechanistic models

In complete mechanistic models there is a correspondence between the variables in the model and a set of observables of the modelled natural system (synaptic connections, metabolic pathways, number of genes, etc.); the model is meant to be realistic, at least to a particular level of mechanistic accuracy. This mechanistic correspondence with the modelled object is exploited to discover which variables and parameter values are crucial to achieve a particular behaviour or functionality, to manipulate the model in ways not accessible to the manipulation of the modelled system and, once the model is adjusted to its object, it might even be used for predictive purposes. Realistic mechanistic models need to integrate many different contingencies and parallel mechanisms that

altogether contribute to the production of the phenomena under study. Most of these models are generally incorporated to existing scientific fields and research programs (microbiology, ethology, neuroscience, etc.) and, although inspired in previous ALife models and modelling techniques, they are not commonly seen in ALife meetings or publications due to their high specificity. Some of these models represent the most complete understanding of a given natural systems we can nowadays hope to achieve: maximally simplified (but still complex), manipulable and predictive artefacts that unfold the causal structure of a particular living phenomena. Examples of realistic and mechanistic models can be found at Webb and colleagues' detailed model of a cricket's phonotaxis (Horchler et al 2004) or Shimizu et al's (2004) model of bacterial chemotaxis.

Conclusions

Almost 20 years after its birth, back in 1987s foundational meeting in Los Alamos, ALife is still alive. Its contributions have been many and have spread over different fields (cognitive science, origin of life, evolutionary thinking, linguistics, robot engineering, theoretical biology, etc.). The most important of its contributions has been the toolkit of simulation and modelling techniques that ALife has brought forth; accompanied by the conceptual and theoretical transformations that these tools have permitted.

The ALife research program has produced a new way to do science, where computers are used to explore the implications of conceptual theories and models, largely in the absence of any direct empirical evaluation. In order to avoid a disappointing itineracy within a playground of virtual-toy-worlds, ALife Models need to be carefully designed, contextualized and manipulated within an epistemic framework that explicitly addresses how and with which purpose the model should be evaluated.

A detailed methodological handbook to improve the epistemic efficiency of ALMs is far beyond the scope of this paper. Nonetheless a set of preliminary provisos might be worth taking home: a) to make explicit which assumptions are built-in into the simulation (as abstractions or simplifications) and which belong to its framework of interpretation, b) to state in advance the purpose and specific hypothesis that the model is aimed to test (including how would it be possible for it to fail) and c) to declare the scope of the model and the relationship it is mean to bear with existing theories or empirical data.

References

Ackley, D. and Littman, M. (1992). Interactions between Learning and Evolution. In Langton, et al. editors, *Proc. Artificial Life II*, Redwood City CA: Addison-Wesley, pp. 487—509.

- Barabási, A. (2002). *Linked: How Everything is Connected to Everything Else*. Perseus Publishing.
- Barabási, A. and Reka, A. (1999). Emergence of scaling in random networks. *Science*, 286:509—512.
- Bechtel, W. & Richardson, R. (1993). *Discovering Complexity*. Princeton University Press.
- Bedau, M. A. (1998). Philosophical content and method of artificial life. In Bynum, T.W. and Moor, J.H. editors. *The digital phoenix: How computers are changing philosophy*. Oxford, UK: Basil Blackwell. pp. 135—152 .
- Beer, R.D. (2003). The dynamics of active categorical perception in an evolved model agent. *Adaptive Behavior*, 11(4): 209—243.
- Bonabeau, E.W. and Theraulaz, G. (1994). Why do we need artificial life? *Artificial Life* 1: 303—325.
- Casti, J. (1997). *Would-be Worlds*. New York: John Wiley.
- Deneubourg, J.L., Theraulaz, G. and Beckers, R. (1991). Swarm-Made Architectures. In Varela, F. and Bourgine, P. editors. *Proc. ECAL 1991*. Cambridge, MA: MIT Press/Bradford Books, pp. 123—133.
- Dennett, D. (1994). Artificial Life as Philosophy. *Artificial Life* 1(3): 291—292.
- Di Paolo, E., Noble, J. & Bullock, S. (2000). Simulation models as opaque thought experiments. In Bedau, et al. editors, *Proc. Artificial Life VII*. Cambridge MA: MIT Press, pp. 497—506.
- Emmeche, C. (1994) *The Garden in the Machine*. Princeton University Press.
- Fox Keller, E. (2005). The Century Beyond the Gene. *J. Bioscience* 30(1): 3—10.
- Gershenson, C. (2004). Updating Schemes in Random Boolean Networks: Do They Really Matter? In Pollack, et al. editors, *Proc. Artificial Life IX*. Cambridge MA: MIT Press, pp. 238—243.
- Hinton G. and Nowlan, S. (1987). How Learning can Guide Evolution, *Complex Systems* 1: 495—502.
- Horchler, A., Reeve, R., Webb, B. and Quinn, R. (2004). Robot phonotaxis in the wild: a biologically inspired approach to outdoor sound localization. *Advanced Robotics*, 18(8):801—816.
- Horgan, J. (1995). From Complexity to Perplexity. *Scientific American* 272:104—109.
- Kauffman, S. (1969). Metabolic Stability and Epigenesis in Randomly Constructed Genetic Nets. *Journal of Theoretical Biology*, 22: 437—467.
- Kauffman, S. (1993). *The origins of order*. Oxford: Oxford University Press.
- Langton, C. (1989). Artificial Life. In Langton, C. editor, *Artificial Life*. Redwood City CA: Addison-Wesley, pp. 1—47.
- Langton, C. (1994) Editor's Introduction. *Artificial Life* 1(1/2): v—viii.
- Liviu, A. and Luke, S. (2004). Ant Foraging Revisited. *Proc. of Artificial Life IX*, pp. 569—574.
- Moreno, A. (2000). Artificial Life as a bridge between Science and Philosophy. In Bedau, M.A., et al. editors, *Proc. Artificial Life VII*. Cambridge, MA: MIT Press, pp 507—512.
- Pattee, H. (1989). Simulations, Realizations and Theories of Life. In Langton, C. editor. *Artificial Life*. Redwood City, CA: Addison Wesley, pp. 63—77.
- Seth, A.K. (1998). Evolving action selection and selective attention without actions, attention, or selection. In Pfeifer, R., et al. editors, *Proc. of the Fifth International Conference on the Simulation of Adaptive Behavior*, Cambridge, MA: MIT Press, pp. 139—147.
- Shimizu, T.S., Aksenov, S.V. and Bray, D. (2003). A spatially extended stochastic model of the bacterial chemotaxis signalling pathway. *Journal of Molecular Biology*, 329:291—309.
- Sismondo, S. (1999). Models, Simulations, and Their Objects. *Science in Context* 12(2): 247—260 .
- Suzuki, R. and Arita, T. (2004). Drastic Changes in Roles of Learning in the Course of Evolution, *Proc. of Artificial Life IX*, pp. 369—374.
- Vickerstaff, R. J., and Di Paolo, E. A. (2005). An evolved agent performing efficient path integration based homing and search. *Proc. ECAL 2005*, Springer Verlag, pp. 221—230.
- Webb, B. (2001). Can robots make good models of biological behaviour? *Behavioral and Brain Sciences*, 24:1033—1105.
- Wheeler, M., Bullock, S., Di Paolo, E., Noble, J., Bedau, M., Husbands, P., Kirby, S. and Seth, A. (2002). The view from elsewhere: perspectives on ALife modelling, *Artificial Life* 8: 87—100.

Acknowledgements

Xabier Barandiaran holds a doctoral fellowship from the Basque Government BFI03371-AE. Funding was also provided by the Comisión Interministerial de Ciencia y Tecnología (CICYT) of Spain, Grant BFF 2002-03294, and BMC2003-06957, and by the University of the Basque Country (UPV/EHU) Grant 9/upv 00003.230-15840/2004. Special thanks to the discussion group (Jon Umerez, Jesús Siqueros and Antonio Casado) of the University of the Basque Country where a draft of this paper was first presented.